

COL 380 : A1

Sidharth Agarwal - 2019CS50661

January 16, 2022

Profiling:

So for profiling I majorly focused on valgrind. After compiling the classify.cpp, run the following commands in the order given for running 1 iteration over 4 threads.

- `valgrind -tool=cachegrind -cachegrind-out-file=out ./classify rfile dfile 1009072 4 1`
- `cg_annotate out /home/sidharth/Parallel380/A1/classify.cpp > profile.txt`

[illegible]

For gprof simply use the command `gprof classify gmon.out ; analysis.txt` and we will get the profiling.

```

Flat profile:

Each sample counts as 0.01 seconds.
 %   cumulative   self           self      total
time  seconds    seconds   calls   Ts/call   Ts/call  name
100.27      8.32      8.32              6     0.00     0.00  readRanges(char const*)
 0.00      8.32      0.00              6     0.00     0.00  classify(Data&, Ranges const&, unsigned int)
 0.00      8.32      0.00              6     0.00     0.00  timedwork(Data&, Ranges const&, unsigned int)
 0.00      8.32      0.00              1     0.00     0.00  _GLOBAL_sub_I_Z8classifyR4DataRK6Rangesj
 0.00      8.32      0.00              1     0.00     0.00  _GLOBAL_sub_I_Z9timedworkR4DataRK6Rangesj

%          the percentage of the total running time of the
time       program used by this function.

cumulative a running sum of the number of seconds accounted
seconds    for by this function and those listed above it.

```

Initial Analysis:

So we carefully look at the profile files of this program, we can easily see that the majority of cache misses occur in inner for loops of both the parallelised code sections. More specifically these four lines of codes *int v = D.data[i].value = R.range(D.data[i].key), if(D.data[d].value == r), D2.data[rangecount[r-1]+rcount++] = D.data[d];*.

Approaches Tested:

rFactor:

So now we can see that there is rangecounter array which is present in the 2nd parallel loop, this results in a lot of cache misses. So a probable reason for that could be false memory sharing. Since the same rangecount chunk of array gets loaded in each of the threads, so we can instead add padding to this array, to avoid false sharing. Now unfortunately this approach doesn't lead to a lot of reduction in time and is only able to reduce time from 430ms to 400ms on average. Similarly I also tested increasing the size of Counter object from 32 to 1028 and padding the counts array as well and it still didn't give a reduction in time nor reduction in cache misses.

Parallel for

So in first loop we are using omp parallel and assigning the data to each thread based on the modulus of the data item, so I decided to test by using parallel for instead so that all the neighbouring elements of Data array are in the same thread. Like sort of a partition of elements among the threads. Unfortunately the approach resulted in increase in time from around 430 ms to 520ms.

Partition over modulus

As we can currently observe that in the first parallel computation we divide the data across different threads based on the modulus of the data point, now to avoid false sharing memory we could instead partition the whole data into continuous segments and then assign it to different threads. Unfortunately it doesn't also lead to a reduction in time.

Algorithm of second parallel loop:

So we can observe that the second parallelised loop is first traversing over all the elements of Ranges and then tested whether each Data element lies in that range. This is highly inefficient since we already know where exactly each Data element lies. Hence instead we can simply iterate over all the elements of Data and add them to their respective bins(Ranges) will keeping a track of how many total elements eventually will be there in that bin and how many have been added yet. This resulted in a decrease in time from around 430ms to 220ms.

Note: This was not improving parallelism, but still mentioning an approach I worked on.